

Lab Report

Part 1:

Problem 1.

main memory = 2^8 bytes
8 lines.


block size = 4 bytes

(1). tag = ? line # = ? byte # = ? for an 8-bit address?

Since block size = 4 bytes, offset = $2^2 \Rightarrow 2$ bits

To represent 8 lines = $2^3 = 8 \Rightarrow 3$ bits

tag = $8 - 3 - 2 = 3$ bits



$$\frac{2^8}{2^5} = 2^3 = 8$$

(2). Into what line would bytes with each of the following address be stored?

000 110 11 $\Rightarrow 110 = \text{line } 6$

001 101 00 $\Rightarrow 101 = \text{line } 5$

110 100 00 $\Rightarrow 100 = \text{line } 4$

101 010 10 $\Rightarrow 010 = \text{line } 2$

$$\frac{2^5}{2^2} = 2^3 = 8$$

(3). Suppose the byte with address 1010001 is stored in the cache. What are the addresses of the other bytes stored along with it?

000 \rightarrow line 0

the block size = 2 bits

\Rightarrow 1010 0000

1010 0001

1010 0010

1010 0011

} the all bytes stored in the block
1010 0000, 1010 0010 and 1010 0011 are the other 3 bytes.

(4) How many total bytes of memory can be stored in the cache?
cache size = $(2^3 \text{ lines}) \times (2^2 \text{ bytes per line}) = 2^5 = 32 \text{ bytes}$

(5) Why is tag also stored in the cache?

To distinguish between 2^3 different blocks that can be stored in the same cache line.

Problem 2: following code:

```
cout << "Hello World"; ①
cin >> a; ②
for (i=0; i<50; i++) ③
    cout << i; ④
```

(1). Give one example of the spatial locality in the code?

spatial locality: instants near recently executed instants likely to be executed soon.

In the for loop, $a[i]$ and $a[i+1]$ will be next to each other in the memory, so if $a[i]$ is used, the next iteration would likely to be $a[i+1]$. And it is executed sequentially.

(2). Give one example of the temporal locality in the code?

Temporal locality: recently executed instants would likely to be executed again; inner loops.

cout is executed first in line 1, it is also used in line 4, which would be likely to be executed again first than cin .

Part 3:

Part 3:

$A[0] \rightarrow \text{address } 0 \times 0$

$B[0] \rightarrow \text{address } 0 \times 100$

num. elements = 8, 4-way associative

cache size = 256

block size = 16

(2) offset number = $\log_2(\text{block size}) \Rightarrow 2^4 = 16 \Rightarrow 4 \text{ bits}$

set number = $\log_2\left(\frac{\text{cache size}}{\text{block size} \times \text{associative}}\right)$
 $= \log_2\left(\frac{256}{16 \times 4}\right) = 4 = 2^2 \Rightarrow 2 \text{ bits}$

tag = $26 - 4 - 2 = 20 \text{ bits}$

addr. $0 \times 0 \Rightarrow \dots 0000 \dots 00 \dots 00 \dots 00$ } in binary
 $0 \times 100 \Rightarrow \dots 0100 \dots 00 \dots 00 \dots 00$
 tag set offset

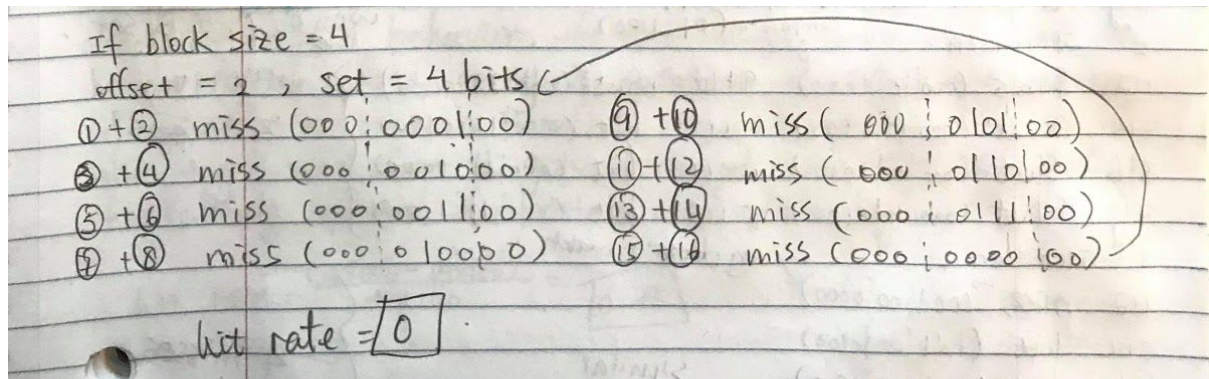
In memory:

$a[0]$	$a[1]$	$a[2]$	$a[3]$	$a[4]$	$a[5]$	$a[6]$	$a[7]$
0×0	0×4	0×8	$0 \times C$	0×10	0×14	0×18	$0 \times 1C$
$b[0]$	$b[1]$	$b[2]$	$b[3]$	$b[4]$	$b[5]$	$b[6]$	$b[7]$
0×100	0×104	0×108	$0 \times 10C$	0×110	0×114	0×118	$0 \times 11C$

① miss $\rightarrow \text{tag} = 0$ ①+② miss (1111)
 ③ miss $\rightarrow \text{tag} = 4$ ①+② hit (10100)
 ③+④ hit $\rightarrow \text{tag} = 0$ ③+④ hit (10000)
 ⑤+⑥ hit $\rightarrow \text{tag} = 4$ ⑤+⑥ hit (11100)
 ⑦+⑧ hit $\rightarrow \text{tag} = 0/4/11/06$ Set 1
 hit rate (Block size = 16) = $12 \div 16 = 0.75$

if block size = 8 : offset bit = 3 bits / set bits = 3 bits / tag = 4 bits

①+③ miss (000000) ①+④ miss (010000)
 ③+④ hit (000100) ①+③ hit (010100) hit rate = $\frac{8}{16} = 0.5$
 ③+⑥ miss (001000) ③+⑧ miss (011000)
 ⑦+⑧ hit (001100) ⑤+⑥ hit (011100)



The results between hand crafted calculation and simulator is the same. As the block size gets smaller, the offset + set digits varies, and when block size is smaller, as address changes the rate of miss is much higher compared to bigger cache block size since now it has higher chance to enter different sets.

Part 4:

Part 4: Both spatial & temporal locality

(2)

num. element = 8 / 4-way associative / cache size = 256
 line size = 16
 offset # = $\log_2(\text{line size}) = 2^4 = 16 \Rightarrow 4 \text{ bits}$
 set # = $\log_2\left(\frac{\text{cache size}}{\text{line size} \times \text{associative}}\right) = \frac{256}{16 \times 4} = 4 = 2^2 \Rightarrow 2 \text{ bits}$
 tag # = $24 - 4 - 2 = 18 \text{ bits}$

a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]	a[8]	a[9]
0	4	8	C	10	14	18	1C	20	24
b[0]	b[1]	b[2]	b[3]	b[4]	b[5]	b[6]	b[7]	b[8]	b[9]
100	104	108	10C	110	114	118	11C	120	124

When iteration = 1:

1+2: miss (000:000)	9+10: miss (01:0000)
3+4: hit (00:0100)	11+12: hit (01:0100)
5+6: hit (00:1000)	13+14: hit (01:1000)
7+8: hit (00:1100)	15+16: hit (01:1100)

hit rate = $\frac{12}{16} = 0.75$

When iteration = 5:

first iteration: 12 hits, 4 misses
 2nd iteration: compare to (01:1100)

1+2: miss (000:10:0000)	9+10: miss (000:11:0000)
3+4: hit (000:10:0100)	11+12: hit (000:11:0100)
5+6: hit (000:10:1000)	13+14: hit (000:11:1000)
7+8: hit (000:10:1100)	15+16: hit (000:11:1100)

3rd iteration: tag doesn't match

1+2: miss (001:00:0000)	
3+4: hit (001:00:0100)	
5+6: hit (001:00:1000)	similar
7+8: hit (001:00:1100)	

4th iteration:

4 misses

Diagram: 4-way set, 2 bits set, 4 misses

$1+2$: miss (001:0:0000)
 $3+4$: hit (001:10:0100)
 \dots
 $9+10$: miss (001:11:0000)
 $11+12$: hit (001:11:0100) } 4 misses

5^{th} iteration =
 $1+2$: miss (010:00:0000)
 $3+4$: hit (010:00:0100)
 \dots
 $9+10$: miss (010:01:0000)
 $11+12$: hit (010:01:0100) } 4 misses

hit rate = $\frac{60}{80} = \boxed{0.75}$

when iteration = 10
 6^{th} iteration: (now tag in set 0 is 010)
 $1+2$: miss (010:10:0000)
 $3+4$: hit (010:10:0100)
 \dots
 $9+10$: miss (011:11:0000)
 $11+12$: hit (011:11:0100)

Overall hit rate (iter = 10) = $\frac{120}{160} = \boxed{0.75}$

100 Iteration:
 according to the behavior, the tag in the cache will start to be overwritten after every 4 iteration, therefore the pattern of misses = 4 in each iteration won't be affected even if the tag & set & offset # gets overlapped, since the LRU policy has already updated the least used data tag.

hit rate = $\frac{(16 \times 100) - (4 \times 100)}{16 \times 100} = \boxed{0.75}$

(5). if use uint_8 instead of int, the address increment is 1 byte instead of 4 byte, therefore each iteration only have 2 misses. The hit rate is much higher than 0.75 ($\frac{14 \times 100 - 2 \times 100}{16 \times 100} = \boxed{0.94}$)

Part 5:


Part 5:

Case 1: 100 iterations, strides = 1

a[0]	a[1]	a[2]	a[3]	a[4]	...	a[127]
0	4	8	12	16	...	1408
a[256]	a[257]					a[511]
200	204	...				3840

Iteration 1:

$$2 \times (128 \div 8) = 2 \times 16 = 32 \text{ misses}$$



for 1 iteration:

1 - every 8 num. element = 2 misses

2 - cache is full,

3 - for 100 iteration:

- tag & set & offset never gets overlapped

$$\text{hit rate} = \frac{(128 \times 100) - (32 \times 100)}{128 \times 100} = \boxed{0.75}$$

same applies to 10 iterations & 50 iterations.

Case 2: 100 iterations, strides = 2

a[0]	a[2]	a[4]	a[6]	a[8]	...
0	16	32	48	64	...


1: miss (00:00:0000) 5: miss (00:10:0000)

2: hit (00:00:1000) 6: hit (00:10:1000)

3: miss (00:01:0000) 7: miss (00:11:0000)

4: hit (00:01:1000) 8: hit (00:11:1000)

9: miss (01:00:0000)



65: miss (1000:00:0000)

each iteration: 4 misses $\times 16 = 64$ misses

$$100 \text{ iterations: } \frac{(128 \times 100) - (64 \times 100)}{128 \times 100} = \boxed{0.5}$$

- LRU cache table starts overwriting every 64 elements.

strides = 8

a[0]	a[8]	a[16]	a[24]	a[32]	...
0	20	40	60	80	...

① miss (000:00:0000) ⑤ miss (60:00:0000)
 ② miss (000:10:0000) ⑥ miss (60:10:0000)
 ③ miss (001:00:0000) ⑦ miss (011:00:0000)
 ④ miss (001:10:0000) ⑧ miss (011:10:0000)
 ...

hit rate = 0 for 100 iterations, as well as 10/50 iterations...

strides = 4

a[0]	a[4]	a[8]	a[16]	a[20]	...
0	10	20	30	40	...

① miss (000:00:0000) ⑤ miss (001:00:0000)
 ② miss (000:01:0000)
 ③ miss (000:02:0000)
 ④ miss (000:03:0000)

hit rate = 0 for 100 iterations, as well as 10/50 iterations

case 16:

a[0]	a[16]	a[32]	...
0	40	80	...

hit rate = 0 !!!
 so inefficient !!!

Since here it goes spatial first, then temporal very bad.

strides number is a very inefficient way, it wastes cache memory as strides

case 21 & 64 = even worse ...

Part 6:

Case 1: direct mapped

offset digit = $\log_2(\text{line_size}) = 4$ bits

set digit = $\log_2(256/(16 \times 1)) = 4$ bits

tag digit = $24 - 4 - 4 = 16$ bits

number of sets = $256 / (16 \times 1) = 16$

1: miss 000 0000 0000

2: hit 000 0000 0100

3: hit 000 0000 1000

```

4:  hit 000 0000 1100
5:  miss 000 0001 0000
6:  hit 000 0001 0100
7:  hit 000 0001 1000
8:  hit 000 0001 1100
9:  miss 000 0010 0000
10: hit 000 0010 0100
11: hit 000 0010 1000
12: hit 000 0010 1100
13: miss 000 0011 0000
14: hit 000 0011 0100
15: hit 000 0011 1000
16: hit 000 0011 1100

```

...

For every 8 elements, 2 misses happens, for every 128 elements, 32 misses happens,
therefore:

$$\text{hit rate} = \frac{(128 \times 100 - (32 \times 100))}{128 \times 100} = 0.75$$

The final cache content would all be tag = 199.

Case 2: 2-way associative

offset digit = $\log_2(\text{line_size}) = 4$ bits

set digit = $\log_2(256 / (16 \times 2)) = 3$ bits

tag digit = $24 - 3 - 4 = 17$ bits

number of sets = $256 / (16 * 2) = 8$

```

1:  miss 0000 000 0000
2:  hit 0000 000 0100
3:  hit 0000 000 1000
4:  hit 0000 000 1100
5:  miss 0000 001 0000
6:  hit 0000 001 0100
7:  hit 0000 001 1000
8:  hit 0000 001 1100
9:  miss 0000 010 0000
10: hit 0000 010 0100
11: hit 0000 010 1000
12: hit 0000 010 1100
13: miss 0000 011 0000
14: hit 0000 011 0100
15: hit 0000 011 1000
16: hit 0000 011 1100
17: miss 0000 100 0000

```

...

For every 8 elements, 2 misses happens, for every 128 elements, 32 misses happens, therefore:

$$\text{hit rate} = \frac{(128 \times 100 - (32 \times 100))}{128 \times 100} = 0.75$$

The final cache content would all be tag = 399 (00 01100 01111).

Case 3: 4-way associative

offset digit = $\log_2(\text{line_size}) = 4$ bits

set digit = $\log_2(256/(16 \times 4)) = 2$ bits

tag digit = $24 - 2 - 4 = 18$ bits

number of sets = $256 / (16 * 4) = 4$

```

1: miss 00000 00 0000
2:  hit 00000 00 0100
3:  hit 00000 00 1000
4:  hit 00000 00 1100
5: miss 00000 01 0000
6:  hit 00000 01 0100
7:  hit 00000 01 1000
8:  hit 00000 01 1100
9: miss 00000 10 0000
10: hit 00000 10 0100
11: hit 00000 10 1000
12: hit 00000 10 1100
13: miss 00000 11 0000
14: hit 00000 11 0100
15: hit 00000 11 1000
16: hit 00000 11 1100
17: miss 00001 00 0000 (different tag)

```

...

For every 8 elements, 2 misses happens, for every 128 elements, 32 misses happens, therefore:

$$\text{hit rate} = \frac{(128 \times 100 - (32 \times 100))}{128 \times 100} = 0.75$$

The final cache content would all be tag = 799 (00 11000 11111).

Case 4: fully associative

offset digit = $\log_2(\text{line_size}) = 4$ bits

set digit = $\log_2(256/(16 \times 16)) = 0$ bits

tag digit = $24 - 2 - 4 = 18$ bits

number of sets = 1

```

1: miss 00000000 0000
2:  hit 00000000 0100
3:  hit 00000000 1000
4:  hit 00000000 1100
5: miss 00000001 0000

```



```
6: hit 0000001 0100
7: hit 0000001 1000
8: hit 0000001 1100
9: miss 0000010 0000
10: hit 0000010 0100
11: hit 0000010 1000
12: hit 0000010 1100
13: miss 0000011 0000
14: hit 0000011 0100
15: hit 0000011 1000
16: hit 0000011 1100
17: miss 0000100 0000 (different tag)
```

...

For every 8 elements, 2 misses happens, for every 128 elements, 32 misses happens, therefore:

$$\text{hit rate} = \frac{(128 \times 100 - (32 \times 100))}{128 \times 100} = 0.75$$

The final cache content would all be tag = 3199 (11 00011 1111).